

# Intro to R for Data Science

Beginner's workshop

AbdulMajedRaja RS

# About Me


- Studied at **Government College of Technology, Coimbatore**
- Bengaluru R user group **Organizer**
- R Packages Developer ( `coinmarketcapr`, `itunesr` )

# Disclaimer:

- This workshop is **NOT** going to make you a Data Scientist **in a day**.
- The objective is to help you get a flavor of R and how it is used in Data Science
- Thus, get you ready to embark on your own journey to become a Data Scientist who uses R

# Content:

This presentation's content is heavily borrowed from the book **R for Data Science** by **Garrett Golemund** and **Hadley Wickham**



# About R

- R is a language and environment for statistical computing and graphics.  
(Ref: [r-project.org](https://r-project.org))
- R was created by Ross Ihaka and Robert Gentleman at the University of Auckland, New Zealand
- R is Free
- R can be extended (easily) via `packages`.
- R is an interpreted language



# R Interpreter / Console / GUI

Demo

# About RStudio

- RStudio is a **free and open-source IDE** for R, released by the company **RStudio, Inc.**
- RStudio and its team regularly contribute to R community by releasing new packages, such as:
  - tidyverse
  - shiny
  - knitr



# RStudio


## Demo



# R Ecosystem

Like Python, R's strength lies in its Ecosystem. **Why R?** - R Packages

## Growth



Source: @daroczig

# Basics of R Programming

# Hello, World!

The traditional first step - **Hello, World!**:

# Hello, World!

The traditional first step - **Hello, World!**:

```
print("Hello, World!")
```

```
## [1] "Hello, World!"
```

That's one small step for a man, one giant leap for mankind

Neil Armstrong

# Arithmetic Operations

```
2 + 3
```

```
## [1] 5
```

```
50000 * 42222
```

```
## [1] 2111100000
```

```
2304 / 233
```

```
## [1] 9.888412
```

```
(33 + 44) * 232 / 12
```

```
## [1] 1488.667
```

# Assignment Operators

**<- Arrow (Less-than < and Minus -)**

**= (Equal Sign)**

```
(x <- 2 + 3)
```

```
## [1] 5
```

```
(y = x ** 4)
```

```
## [1] 625
```

```
5 * 9 -> a  
a + 3
```

```
## [1] 48
```

# Objects

- The entities R operates on are technically known as objects.

Example: Vector of numeric

```
vector_of_numeric <- c(2,4,5)
```

```
typeof(vector_of_numeric)
```

```
## [1] "double"
```

# Vectors

- Atomic Vectors - Homogeneous Data Type
  - logical
  - integer
  - double
  - character
  - 
  -
- Lists - (Recursive Vectors) Heterogeneous Data Type
- NULL is used to represent absence of a vector

**Vectors + Attributes (Additional Meta Data) = Augmented vectors**

- Factors are built on top of integer vectors.
- Dates and date-times are built on top of numeric vectors.
- Data frames and tibbles are built on top of lists.



# Numeric Vector

```
nummy <- c(2,3,4)
nummy_int <- c(1L,2L,3L)
```

```
typeof(nummy)
```

```
## [1] "double"
```

```
typeof(nummy_int)
```

```
## [1] "integer"
```

```
is.numeric(nummy)
```

```
## [1] TRUE
```

```
is.numeric(nummy_int)
```

```
## [1] TRUE
```

```
is.double(nummy)
```

# Character Vector

```
types <- c("int", "double", "character")
```

```
types
```

```
## [1] "int"      "double"    "character"
```

```
typeof(types)
```

```
## [1] "character"
```

```
length(types)
```

```
## [1] 3
```

```
is.numeric(types)
```

```
## [1] FALSE
```

```
is.character(types)
```

```
## [1] TRUE
```

# Logical Vector

```
logicals <- c(TRUE,F,TRUE,T, FALSE)
```

```
logicals
```

```
## [1] TRUE FALSE TRUE TRUE FALSE
```

# Coersion

## Typcasting - Explicit

```
money_in_chars <- c("20", "35", "33")  
typeof(money_in_chars)
```

```
## [1] "character"
```

```
money_money <- as.numeric(money_in_chars)  
money_money
```

```
## [1] 20 35 33
```

```
typeof(money_money)
```

```
## [1] "double"
```

# Typecasting - Implicit

```
money_money <- as.numeric(money_in_chars)
```

```
money_money
```

```
## [1] 20 35 33
```

```
typeof(money_money)
```

```
## [1] "double"
```

```
new_money <- c(money_money, "33")
```

```
new_money
```

```
## [1] "20" "35" "33" "33"
```

```
typeof(new_money)
```

```
## [1] "character"
```

# Vector - Accessing

```
month.abb #in-built character vector with Month Abbreviations
```

```
## [1] "Jan" "Feb" "Mar" "Apr" "May" "Jun" "Jul" "Aug" "Sep" "Oct" "Nov"  
## [12] "Dec"
```

```
month.abb[2]
```

```
## [1] "Feb"
```

```
month.abb[4:7]
```

```
## [1] "Apr" "May" "Jun" "Jul"
```

```
month.abb[c(2,5,7,10)]
```

```
## [1] "Feb" "May" "Jul" "Oct"
```

# Vector Manipulation

## Appending

```
days <- c("Mon", "Tue", "Wed")
```

```
days
```

```
## [1] "Mon" "Tue" "Wed"
```

```
week_end <- c("Sat", "Sun")
```

```
more_days <- c(days, "Thu", "Fri", week_end)
```

```
more_days
```

```
## [1] "Mon" "Tue" "Wed" "Thu" "Fri" "Sat" "Sun"
```

# Vector - Arithmetic

```
set.seed(122)

so_many_numbers <- runif(10, min = 10, max = 100)

so_many_numbers
```

```
## [1] 91.45185 91.61657 27.14995 13.68211 62.11661 66.31451 71.14383
## [8] 10.25104 13.21914 63.69918
```

```
so_many_numbers * 200
```

```
## [1] 18290.370 18323.314 5429.989 2736.422 12423.322 13262.902 14228.767
## [8] 2050.209 2643.828 12739.836
```



# Factors

- In R, factors are used to work with categorical variables, variables that have a fixed and known set of possible values.
- Useful with Characters where non-Alphabetical Ordering is required

```
days <- c("Thu", "Wed", "Sun")  
  
sort(days)
```

```
## [1] "Sun" "Thu" "Wed"
```

```
week_levels <- c("Mon", "Tue", "Wed", "Thu", "Fri", "Sat", "Sun")  
  
(days_f <- factor(days, levels = week_levels))
```

```
## [1] Thu Wed Sun  
## Levels: Mon Tue Wed Thu Fri Sat Sun
```

```
sort(days_f)
```

# List

Lists are a step up in complexity from atomic vectors: each element can be any type, not just vectors.

```
(a_list <- list("abcd",123,1:12,month.abb))
```

```
## [[1]]  
## [1] "abcd"  
##  
## [[2]]  
## [1] 123  
##  
## [[3]]  
## [1] 1 2 3 4 5 6 7 8 9 10 11 12  
##  
## [[4]]  
## [1] "Jan" "Feb" "Mar" "Apr" "May" "Jun" "Jul" "Aug" "Sep" "Oct" "Nov"  
## [12] "Dec"
```

# List Accessing

```
a_list[[1]]
```

```
## [1] "abcd"
```

```
a_list[[4]][4]
```

```
## [1] "Apr"
```

# Matrix

```
new_m <- matrix(data = 1:12, nrow = 3)
```

```
new_m
```

```
##      [,1] [,2] [,3] [,4]
## [1,]    1    4    7   10
## [2,]    2    5    8   11
## [3,]    3    6    9   12
```

```
new_m * 20
```

```
##      [,1] [,2] [,3] [,4]
## [1,]   20   80  140  200
## [2,]   40  100  160  220
## [3,]   60  120  180  240
```

```
dim(new_m)
```

```
## [1] 3 4
```

# Dataframe

## Tabular Structure

- dimension
- row.names
- col.names

```
colleges <- c("CIT", "GCT", "PSG")  
year <- c(2019, 2018, 2017)  
db <- data.frame(college_names = colleges, year_since = year)  
db
```

```
##   college_names year_since  
## 1           CIT      2019  
## 2           GCT      2018  
## 3           PSG      2017
```

# Dataframe Manipulation

```
db$college_names
```

```
## [1] CIT GCT PSG  
## Levels: CIT GCT PSG
```

```
db[2,2] <- 1990
```

```
db
```

```
##   college_names year_since  
## 1           CIT      2019  
## 2           GCT      1990  
## 3           PSG      2017
```

```
db[, "year_since"]
```

```
## [1] 2019 1990 2017
```

# Loops & Iterators

## For Loop

```
(month_name      month.abb[1:4]) {  
  print(paste("This month", month_name, "beautiful!!!"))  
}
```

```
## [1] "This month Jan beautiful!!!"  
## [1] "This month Feb beautiful!!!"  
## [1] "This month Mar beautiful!!!"  
## [1] "This month Apr beautiful!!!"
```

As you move forward, Check the family of apply functions - `sapply()`, `tapply()`, `lapply()`, `apply()`.

For advanced functional programming, refer `purrr` package

# Logical Operations

## %in% operator

```
iris$Species % % "virginica"
```

```
## [1] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [12] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [23] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [34] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [45] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [56] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [67] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [78] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [89] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [100] FALSE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
## [111] TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
## [122] TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
## [133] TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
## [144] TRUE TRUE TRUE TRUE TRUE TRUE TRUE
```



# Logical Operators

```
1:10 > 5
```

```
## [1] FALSE FALSE FALSE FALSE FALSE TRUE TRUE TRUE TRUE TRUE
```

```
1:10 == 4
```

```
## [1] FALSE FALSE FALSE TRUE FALSE FALSE FALSE FALSE FALSE FALSE
```

```
!1:10 == 4
```

```
## [1] TRUE TRUE TRUE FALSE TRUE TRUE TRUE TRUE TRUE TRUE
```

# Conditions

```
(iris$Sepal.Length[2]>5) {  
  print("it is gt 5")  
}  
  print("it is not")
```

```
## [1] "it is not"
```

```
(iris$Sepal.Length>10) {print("hello")}
```

```
## Warning in if (iris$Sepal.Length > 10) {: the condition has length > 1 and  
## only the first element will be used
```

```
ifelse(iris$Sepal.Length>6, "more_than_10", "les_than_10")
```

```
##      [1] "les_than_10" "les_than_10" "les_than_10" "les_than_10"  
##      [5] "les_than_10" "les_than_10" "les_than_10" "les_than_10"  
##      [9] "les_than_10" "les_than_10" "les_than_10" "les_than_10"  
##     [13] "les_than_10" "les_than_10" "les_than_10" "les_than_10"  
##     [17] "les_than_10" "les_than_10" "les_than_10" "les_than_10"  
##     [21] "les_than_10" "les_than_10" "les_than_10" "les_than_10"  
##     [25] "les_than_10" "les_than_10" "les_than_10" "les_than_10"
```

# Functions

## Types

- Base-R functions (mean(), plot(), lm())
- Package functions (dplyr::mutate(), stringr::str\_detect())
- User-defined functions

```
workshop_hate_message <- (name = "No one", n = 3) {  
  text_to_print <- paste(name, "hate(s)", "this workshop")  
  (i = 1:n) {  
    print(text_to_print)  
  }  
}  
  
workshop_hate_message("All of us", 4)
```

```
## [1] "All of us hate(s) this workshop"  
## [1] "All of us hate(s) this workshop"  
## [1] "All of us hate(s) this workshop"  
## [1] "All of us hate(s) this workshop"
```

# Packages

## Package Installation & Loading

From CRAN (usually Stable Version)

```
install.packages("itunesr")
```

**From Github (usually Development Version)**

```
#install.packages("devtools")  
devtools::install_github("amrrs/itunesr")
```

## Loading

```
(itunesr)
```

# Help

using `help()`

```
help("runif")
```

using ?

```
?sample
```

# Help - Example

```
example("for")
```

```
##
## for> for(i in 1:5) print(1:i)
## [1] 1
## [1] 1 2
## [1] 1 2 3
## [1] 1 2 3 4
## [1] 1 2 3 4 5
##
## for> for(n in c(2,5,10,20,50)) {
## for+   x <- stats::rnorm(n)
## for+   cat(n, ": ", sum(x^2), "\n", sep = "")
## for+ }
## 2: 2.188171
## 5: 1.936692
## 10: 15.34038
## 20: 25.59841
## 50: 49.75875
##
## for> f <- factor(sample(letters[1:5], 10, replace = TRUE))
##
```

# Packages Vignette

```
vignette("dplyr")
```

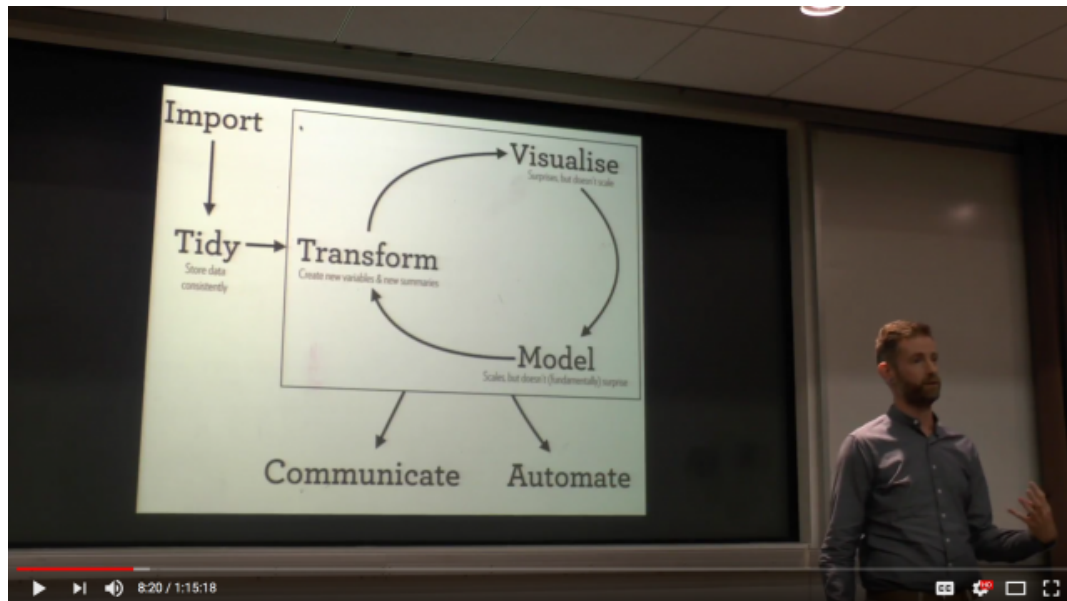
```
browseVignettes("dplyr")
```

# Data wrangling and Visualization using Tidyverse



# Data Science Framework

There are now like, you know, a billion venn diagrams showing you what data science is. But to me I think the definition is pretty simple. Whenever you're struggling with data, trying to understand what's going on with data, whenever you're trying to turn that **raw data into insight and understanding and discoveries**. I think that's **Data Science**." - Hadley Wickham



Source: [Hadley Wickham](#)

# Tidyverse

- An opinionated collection of R packages designed for data science.
- All packages share an underlying design

```
install.packages("tidyverse")
```

## tidyverse packages

```
tidyverse::tidyverse_packages()
```

```
## [1] "broom"      "cli"        "crayon"     "dplyr"      "dbplyr"
## [6] "forcats"    "ggplot2"    "haven"      "hms"        "httr"
## [11] "jsonlite"   "lubridate"  "magrittr"   "modelr"     "purrr"
## [16] "readr"      "readxl"     "reprex"     "rlang"      "rstudioapi"
## [21] "rvest"      "stringr"    "tibble"     "tidyr"      "xml2"
## [26] "tidyverse"
```

# Loading the Library

```
(tidyverse)
```

```
## — Attaching packages ————— tidyverse
```

```
## ✔ ggplot2 3.1.0      ✔ purrr  0.2.5  
## ✔ tibble  2.0.0      ✔ dplyr   0.7.8  
## ✔ tidyr   0.8.2      ✔ stringr 1.3.1  
## ✔ readr   1.3.1      ✔ forcats 0.3.0
```

```
## Warning: package 'tibble' was built under R version 3.5.2
```

```
## — Conflicts ————— tidyverse
```

```
## ✖ dplyr::filter() masks stats::filter()  
## ✖ dplyr::lag()     masks stats::lag()
```

# Input Data

Reading the dataset

```
#kaggle <- read_csv("data/kaggle_survey_2018.csv")  
kaggle <- read_csv("data/kaggle_survey_2018.csv", skip = 1)
```

# Basic Stats

## Dimension (Rows Column)

```
dim(kaggle)
```

```
## [1] 23859    395
```

```
glimpse(kaggle)
```

```
## Observations: 23,859
```

```
## Variables: 395
```

```
## $ `Duration (in seconds)`
```

```
## $ `What is your gender? - Selected Choice`
```

```
## $ `What is your gender? - Prefer to self-describe - Text`
```

```
## $ `What is your age (# years)?`
```

```
## $ `In which country do you currently reside?`
```

```
## $ `What is the highest level of formal education that you have attained or
```

```
## $ `Which best describes your undergraduate major? - Selected Choice`
```

```
## $ `Select the title most similar to your current role (or most recent titl
```

```
## $ `Select the title most similar to your current role (or most recent titl
```

```
## $ `In what industry is your current employer/contract (or your most recent
```

# Dataset Overview

Demo on RStudio

# Data Questions (Business Problem)

- What's the percentage of Male and Female respondents?
- What are the top 5 countries?

# dyplr verbs

- `mutate()` - adds new variables that are functions of existing variables
- `select()` - picks variables based on their names.
- `filter()` - picks cases based on their values.
- `summarise()` - reduces multiple values down to a single summary.
- `arrange()` - changes the ordering of the rows.



# Introducing %>% Pipe Operator

- The pipe, %>%, comes from the magrittr package by Stefan Milton Bache
- **Output of LHS** is given as the **input (first argument) of RHS**

## Example

```
kaggle %>% dim()
```

```
## [1] 23859 395
```

Although doesn't make much sense to use %>% in this context, Hope it explains the function.

# Percentage of Male and Female

- Column name - What is your gender? - Selected Choice

## Pseudo-code

- group\_by the kaggle dataframe on column What is your gender? - Selected Choice
- count the values
- calculate percentage value from the counts

# % of Male and Female - Group By & Count - Method 1

```
kaggle %>%  
  group_by(`What is your gender? - Selected Choice`) %>%  
  summarise(n = n())
```

```
## # A tibble: 4 x 2  
##   `What is your gender? - Selected Choice`      n  
##   <chr>                                     <int>  
## 1 Female                                     4010  
## 2 Male                                     19430  
## 3 Prefer not to say                         340  
## 4 Prefer to self-describe                    79
```

# % of Male and Female - Group By & Count - Method 2

```
kaggle %>%  
  group_by(`What is your gender? - Selected Choice`) %>%  
  count()
```

```
## # A tibble: 4 x 2  
## # Groups:   What is your gender? - Selected Choice [4]  
##   `What is your gender? - Selected Choice`      n  
##   <chr>                                     <int>  
## 1 Female                                     4010  
## 2 Male                                     19430  
## 3 Prefer not to say                        340  
## 4 Prefer to self-describe                  79
```

# % of Male and Female - Group By & Count - Sorted

```
kaggle %>%  
  group_by(`What is your gender? - Selected Choice`) %>%  
  count() %>%  
  arrange(desc(n))
```

```
## # A tibble: 4 x 2  
## # Groups:   What is your gender? - Selected Choice [4]  
##   `What is your gender? - Selected Choice`      n  
##   <chr>                                     <int>  
## 1 Male                                     19430  
## 2 Female                                    4010  
## 3 Prefer not to say                        340  
## 4 Prefer to self-describe                   79
```

# % of Male and Female - Percentage

```
kaggle %>%  
  group_by(`What is your gender? - Selected Choice`) %>%  
  count() %>%  
  ungroup() %>%  
  mutate(perc = round(n / sum(n),2))
```

```
## # A tibble: 4 x 3  
##   `What is your gender? - Selected Choice`      n  perc  
##   <chr>                                <int> <dbl>  
## 1 Female                                4010  0.17  
## 2 Male                                19430  0.81  
## 3 Prefer not to say                     340  0.01  
## 4 Prefer to self-describe                79  0
```

# % of Male and Female - Nice\_Looking\_Table

```
kaggle %>%  
  group_by(`What is your gender? - Selected Choice`) %>%  
  count() %>%  
  ungroup() %>%  
  mutate(perc = round(n / sum(n),2)) %>%  
  knitr::kable(format = "html")
```

What is your gender? - Selected Choice	n	perc
Female	4010	0.17
Male	19430	0.81
Prefer not to say	340	0.01
Prefer to self-describe	79	0.00

But, Wait!!!

Go Back and See

If you have only Male and Female?



Time for some cleaning

In the form of `filter()`ing

# % of Male and Female - Filtered\_Nice

```
kaggle %>%  
  filter(`What is your gender? - Selected Choice` % % c("Male", "Female"))  
  group_by(`What is your gender? - Selected Choice`) %>%  
  count() %>%  
  ungroup() %>%  
  mutate(perc = round(n / sum(n), 2)) %>%  
  knitr::kable(format = "html")
```

What is your gender? - Selected Choice	n	perc
Female	4010	0.17
Male	19430	0.83

An Awkward column name, isn't it??!

# % of Male and Female - All\_Nice\_Table

```
(scales) #for Percentage Formatting
```

```
##  
## Attaching package: 'scales'  
  
## The following object is masked from 'package:purrr':  
##  
##      discard  
  
## The following object is masked from 'package:readr':  
##  
##      col_factor
```

```
kaggle %>%  
  filter(`What is your gender? - Selected Choice` % % c("Male", "Female"))  
  group_by(`What is your gender? - Selected Choice`) %>%  
  count() %>%  
  ungroup() %>%  
  mutate(perc = round(n / sum(n), 2)) %>%  
  mutate(perc = scales::percent(perc)) %>%  
  rename(Gender = `What is your gender? - Selected Choice`,
```

# Top 5 Countries

- Column name - In which country do you currently reside?

## Pseudo-code

- count number of respondents from each country
- arrange countries in descending order based on their count value
- top 5 in the list is the output

# Top 5 Countries - Code

```
kaggle %>%  
  count(`In which country do you currently reside?`) %>%  
  arrange(desc(n)) %>%  
  top_n(5) %>%  
  knitr::kable(format = "html")
```

## Selecting by n

In which country do you currently reside?	n
United States of America	4716
India	4417
China	1644
Other	1036
Russia	879

Is Other a country name???

# Top 5 Countries

```
kaggle %>%  
  filter(!`In which country do you currently reside?` % % "Other") %  
  count(`In which country do you currently reside?`) %>%  
  rename(Country = `In which country do you currently reside?`) %>%  
  arrange(desc(n)) %>%  
  top_n(5) %>%  
  knitr::kable(format = "html")
```

## Selecting by n

Country	n
United States of America	4716
India	4417
China	1644
Russia	879
Brazil	736



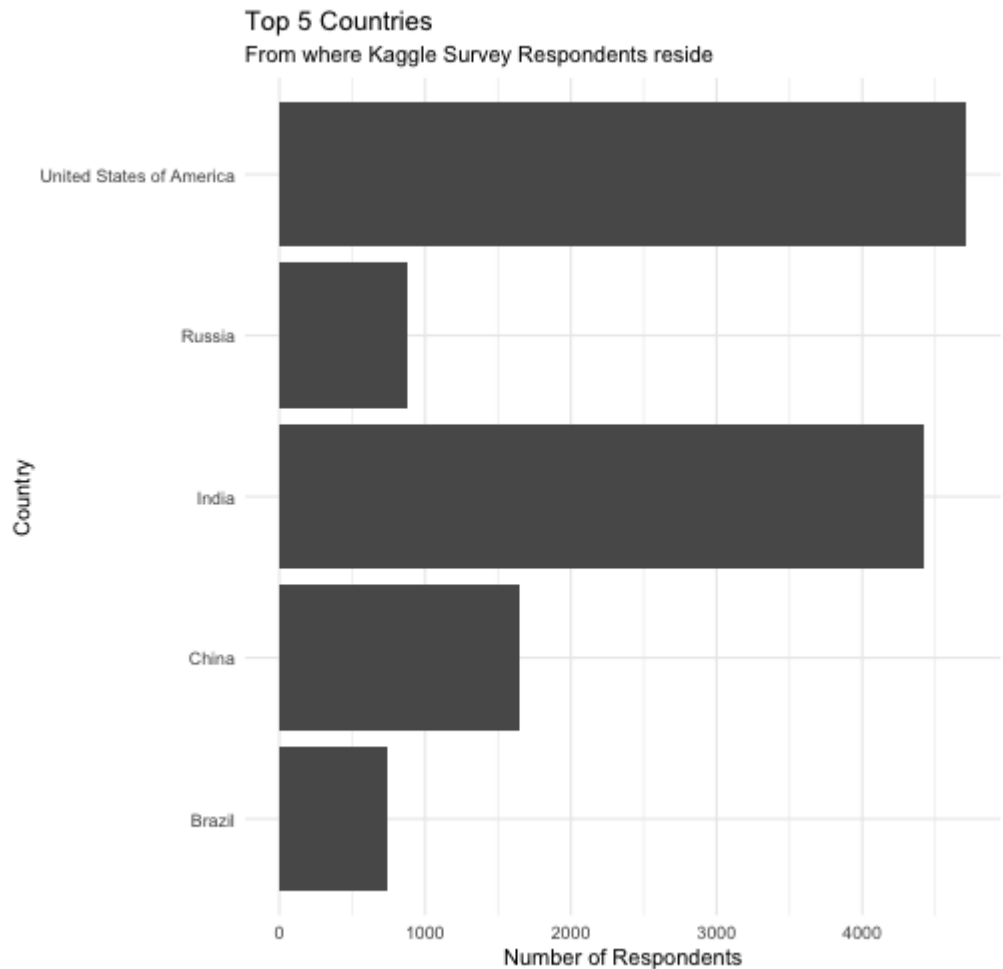
Table is nice, but a visually appealing plot is  
Nicer



# Top 5 Countries - Plot #1

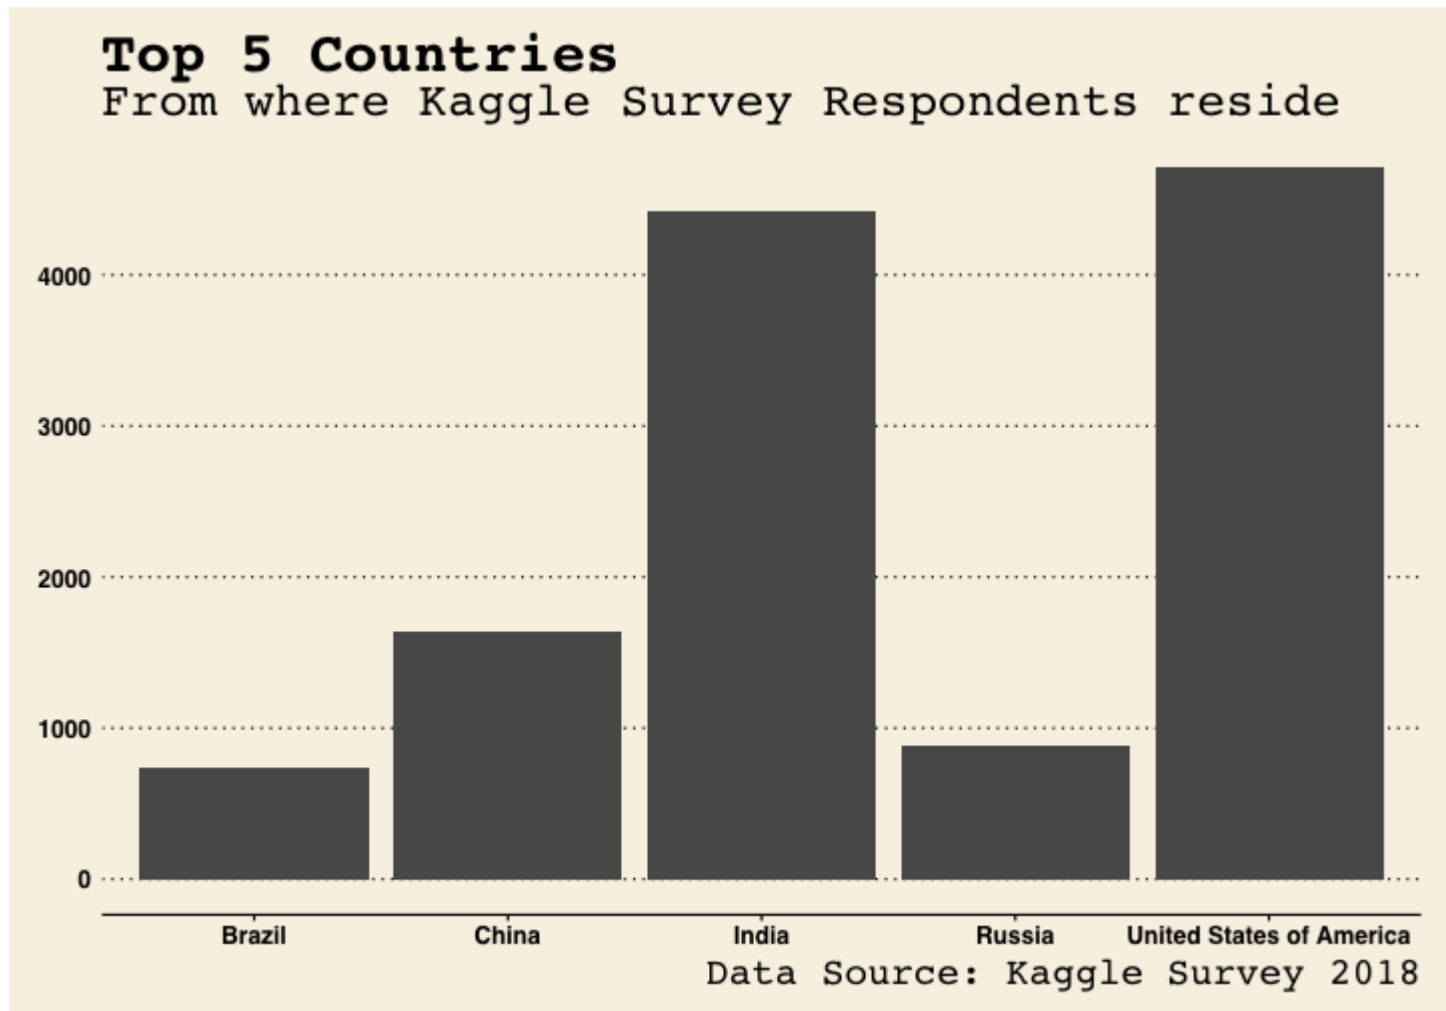
```
kaggle %>%  
  filter(!`In which country do you currently reside?` % % "Other") %  
  count(`In which country do you currently reside?`) %>%  
  rename(Country = `In which country do you currently reside?`) %>%  
  arrange(desc(n)) %>%  
  top_n(5) %>%  
  ggplot() + geom_bar(aes(Country,n), stat = "identity") +  
  coord_flip() +  
  theme_minimal() +  
  labs(title = "Top 5 Countries",  
        subtitle = "From where Kaggle Survey Respondents reside",  
        x = "Country",  
        y = "Number of Respondents",  
        caption = "Data Source: Kaggle Survey 2018")
```

# Top 5 Countries - Plot #2



Data Source: Kaggle Survey 2018

# Top 5 Countries - Plot #3 Themed



# Documentation and Reporting using R Markdown

Demo

# Project Demo

# Object Detection in 3 Lines of R Code

using Tiny YOLO

-Project Demo-

# References

- R for Data Science
- R-Bloggers



# Thanks!

Slides created via the R package **xaringan**.

The chakra comes from **remark.js**, **knitr**, and **R Markdown**.